✓ Congratulations! You passed!

1/1 points

1. Which notation would you use to denote the 3rd layer's activations when the input is the 7th example from the 8th minibatch?		
	$a^{[8]\{7\}(3)}$	
	$a^{[3]\{7\}(8)}$	
	$a^{[8]\{3\}(7)}$	
0	$a^{[3]\{8\}(7)}$	
Corre	ect	
~	1 / 1 points	
2. Which with?	of these statements about mini-batch gradient descent do you agree	
	You should implement mini-batch gradient descent without an explicit for-loop over different mini-batches, so that the algorithm processes all mini-batches at the same time (vectorization).	
	Training one epoch (one pass through the training set) using minibatch gradient descent is faster than training one epoch using batch gradient descent.	
\bigcirc		

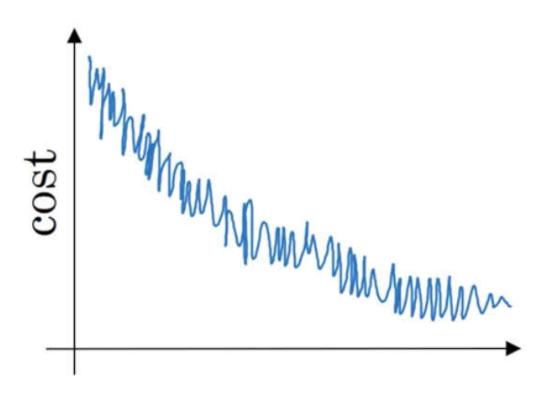
One iteration of mini-batch gradient descent (computing on a single mini-batch) is faster than one iteration of batch gradient descent. Correct 1/1 points 3. Why is the best mini-batch size usually not 1 and not m, but instead something in-between? If the mini-batch size is 1, you end up having to process the entire training set before making any progress. **Un-selected is correct** If the mini-batch size is m, you end up with batch gradient descent, which has to process the whole training set before making progress. Correct If the mini-batch size is 1, you lose the benefits of vectorization across examples in the mini-batch. Correct If the mini-batch size is m, you end up with stochastic gradient descent, which is usually slower than mini-batch gradient descent. **Un-selected is correct**



1/1 points

4.

Suppose your learning algorithm's cost J, plotted as a function of the number of iterations, looks like this:



Which of the following do you agree with?

\bigcirc	Whether you're using batch gradient descent or mini-batch gradient
	descent, something is wrong.

- If you're using mini-batch gradient descent, something is wrong. But if you're using batch gradient descent, this looks acceptable.
- Whether you're using batch gradient descent or mini-batch gradient descent, this looks acceptable.
- If you're using mini-batch gradient descent, this looks acceptable. But if you're using batch gradient descent, something is wrong.

Correct



1/1 points

5.

Suppose the temperature in Casablanca over the first three days of January are the same:

Jan 1st:
$$heta_1=10^oC$$

Jan 2nd: $heta_2 10^o C$

(We used Fahrenheit in lecture, so will use Celsius here in honor of the metric world.)

Say you use an exponentially weighted average with $\beta=0.5$ to track the temperature: $v_0=0$, $v_t=\beta v_{t-1}+(1-\beta)\theta_t$. If v_2 is the value computed after day 2 without bias correction, and $v_2^{corrected}$ is the value you compute with bias correction. What are these values? (You might be able to do this without a calculator, but you don't actually need one. Remember what is bias correction doing.)

$$igcup v_2=10$$
, $v_2^{corrected}=10$

$$v_2=7.5$$
, $v_2^{corrected}=7.5$

$$v_2=7.5$$
 , $v_2^{corrected}=10$

Correct

$$igcup v_2=10$$
, $v_2^{corrected}=7.5$



1/1 points

6.

Which of these is NOT a good learning rate decay scheme? Here, t is the epoch number.

$$\alpha = \frac{1}{\sqrt{t}} \alpha_0$$

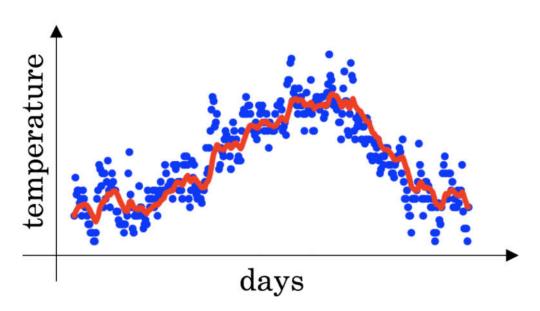
$$igcap lpha = 0.95^t lpha_0$$

$$igcap lpha = e^t lpha_0$$

Correct

$$igcap lpha = rac{1}{1+2*t}\,lpha_0$$

7. You use an exponentially weighted average on the London temperature dataset. You use the following to track the temperature: $v_t = \beta v_{t-1} + (1-\beta)\theta_t.$ The red line below was computed using $\beta = 0.9$. What would happen to your red curve as you vary β ? (Check the two that apply)



Un-selected is correct

Increasing β will shift the red line slightly to the right.

Correct

True, remember that the red line corresponds to $\beta=0.9$. In lecture we had a green line \$\$\beta=0.98\$) that is slightly shifted to the right.

Decreasing β will create more oscillation within the red line.

Correct

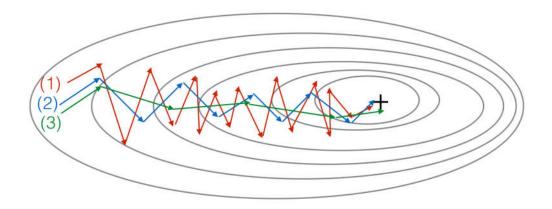
True, remember that the red line corresponds to $\beta=0.9$. In lecture we had a yellow line \$\$\beta=0.98\$ that had a lot of oscillations.

Un-selected is correct



1/1 points

8. Consider this figure:



These plots were generated with gradient descent; with gradient descent with momentum (β = 0.5) and gradient descent with momentum (β = 0.9). Which curve corresponds to which algorithm?

- (1) is gradient descent. (2) is gradient descent with momentum (large β) . (3) is gradient descent with momentum (small β)
- (1) is gradient descent with momentum (small β). (2) is gradient descent. (3) is gradient descent with momentum (large β)
- (1) is gradient descent with momentum (small β), (2) is gradient descent with momentum (small β), (3) is gradient descent
- (1) is gradient descent. (2) is gradient descent with momentum (small β). (3) is gradient descent with momentum (large β)

Correct



1/1 points

Suppose batch gradient descent in a deep network is taking excessively long to find a value of the parameters that achieves a small value for the cost function $\mathcal{J}(W^{[1]},b^{[1]},\ldots,W^{[L]},b^{[L]})$. Which of the following techniques could help find parameter values that attain a small value for \mathcal{J} ? (Check all that apply)

	Try initializing all the weights to zero		
Un-se	Un-selected is correct		
	Try using Adam		
Correct			
	Try better random initialization for the weights		
Correct			
Corre	Try tuning the learning rate $lpha$		
COITE			
	Try mini-batch gradient descent		
Corre	ect		
~	1/1 points		
10. Which	of the following statements about Adam is False?		
\bigcirc	Adam combines the advantages of RMSProp and momentum		
0	Adam should be used with batch gradient computations, not with mini-batches.		
Correct			

We usually use "default" values for the hyperparameters eta_1,eta_2 and $arepsilon$ in Adam ($eta_1=0.9,eta_2=0.999,arepsilon=10^{-8}$)
The learning rate hyperparameter α in Adam usually needs to be tuned.